

# **Designing a Scalable Machine Learning Pipeline for Cyber Threat Detection**

Araynah Dover • CMSE492 • 11/25

# Problem & Motivation

---

Cyber Attacks are increasing and traditional rule-based IDS systems cannot detect new or evolving threats.

Traffic volume is too large for humans to analyze

Detecting anomalies early prevents breaches

# Dataset Overview

---

CICIDS-2017 (all 5 days merged)

2.1 million rows, 105 features

Includes real enterprise traffic + botnet, DoS, brute force, infiltration, XSS, etc.

After cleaning: ≈86 usable features

Target: label\_clean → converted to Binary-Label (0 = benign, 1 = attack)

# **Methodology**


# Why ML Works Here

Traditional IDS limitations:

- Only detects known attacks
- Cannot generalize
- High false-positives

Why ML is better:

- Learns from statistical patterns
- Detects unseen anomalies
- Scales to millions of flows





# Preprocessing Steps



---

- Dropped irrelevant columns (FlowID, timestamp, local columns)
  - Cleaned labels → `label_clean` → `BinaryLabel`
  - Handle extreme skew and outliers (log scaling on heavy-tailed fields)
  - One-hot encoded Protocol
  - Train/test split: Stratified 80/20
  - No SMOTE → use `class_weight='balanced'`
  - `StandardScaler` → fit only on training data
-

# Exploratory Data Analysis



# Exploratory Data Analysis



# Feature Engineering

---

- Dropped redundant + low-variance columns
  - One-hot encoded Protocol
  - Standard scaling
-

# Models Compared

---

## Logistic Regression:

- Linear decision boundary
- Very fast
- Good interpretability
- Uses `class_weight='balanced'`

## Neural Network:

- Can model complex relationships
- Needs regularization
- Harder to interpret

## Random Forest:

- Captures nonlinear patterns
  - Robust to noise
  - Naturally ranks feature importance
-

# Training Strategy

---

- StandardScaler → only on training data
  - class\_weight to address imbalance
  - 5-fold CV for hyperparameters
  - Same train/test split for all models
-

# Evaluation Metrics

---

- Accuracy
  - Precision
  - Recall
  - F1-score
  - ROC-AUC
  - Confusion matrices
-


# Results

# Key Results

	Model	Accuracy	Precision	Recall	F1 Score	ROC-AUC
0	Logistic Regression	0.997000	0.991358	0.996471	0.993908	0.999795
1	Random Forest	0.999883	0.999932	0.999593	0.999762	1.000000
2	Neural Network	0.999700	0.999321	0.999457	0.999389	0.999990

# Model Interpretation

```
git add data/processed/merged.csv  
git commit -m "Add merged dataset"  
git push origin main
```




# Challenges encountered

SMOTE too large → switched to class weight

SHAP too memory-heavy → removed

Large dataset → had to downsample for EDA

Many useless columns → dropped to speed up pipeline



## Future Steps

Try more complex approaches

Try to get a more foundational understanding